

Database Theory

Part I - Topics Covered

- Database Tables - What they are and how they are created
- Data Types - What kind of values can be stored in a column
- Constraints - Enforcing rules on what type of data can go where
- Relationships - How data from one table can reference data from another table

Why learn about databases?

- A website without a database is just a static collection of text
- Databases efficiently store and retrieve information
- This information often comes from HTML form submissions

What is a database?

- Data can be inserted, read, updated, and destroyed from the database
- The database itself lives in a file on your computer
- Databases have one or more **tables**

Tables

- Every table stores some unique category of data
- A **"users" table** might store user information like name, email address, and password.
- A **"products" table** might store information like product name and price.

Tables

+-----+			
users			
+-----+			
id	name	password	
+-----+			
1	Sally	\$2y\$10\$nKvx7rTyK8c10	
2	Ned	\$2y\$10\$DTkFfgtDdRJWV	
3	Sally	\$2y\$10\$mf7C1K72LGx48	
+-----+			

Column

Row/Record

Tables

```
+-----+
| users |
+-----+
| id | name | password |
+-----+
| 1 | Sally | $2y$10$nKvx7rTyK8c10 |
| 2 | Ned | $2y$10$DTkFfgtDdRJWV |
| 3 | Sally | $2y$10$mf7C1K72LGx48 |
+-----+
```

- The ID column is the unique identifier for the particular row it belongs to
- It is referred to as a **PRIMARY KEY**
- Almost every table will include an ID column for storing unique IDs
- The unique ID is an Auto Incrementing ID (automatically assigned)

Tables

```
+-----+
| users |
+-----+
| id | name | password |
+-----+
| 1 | Sally | $2y$10$nKvx7rTyK8c10 |
| 2 | Ned | $2y$10$DTkFfgtDdRJWV |
| 3 | Sally | $2y$10$mf7C1K72LGx48 |
+-----+
```

- You don't need to manually assign the number value of the auto incrementing ID
- We might have more than one person named Sally but these are two different people
- Rows can contain similar data
- A unique ID is used to differentiate rows and create relations (more on relations later)

Tables

```
+-----+
| users |
+-----+
| id | name | password |
+-----+
| 1 | Sally | $2y$10$nKvx7rTyK8c10 |
| 2 | Ned | $2y$10$DTkFfgtDdRJWV |
| 3 | Sally | $2y$10$mf7C1K72LGx48 |
+-----+
```

- The name column contains a simple string value
- The password column contains a "hashed" or encrypted version of the persons actual password

Tables

- The database schema provides its entire structure
- The structure includes **tables** and **columns**
- Every column has a **data type**

Data Types

When defining a table you must specify a data type for every column you add to it. There are many data types and we will not cover all of them. **Some** common data types include:

- **VARCHAR**: Variable length string of characters
- **INTEGER**: Number with no decimal (5)
- **DECIMAL**: Number with decimal (50.10)
- **BOOLEAN**: Stores TRUE/FALSE values (1 or 0)
- **DATETIME**: Stores a date and time value

Please **write these down**. You will need them later.

Database Management Systems

- Tables are created through SQL (Structured Query Language)
- There are many forms of SQL, most of which are fairly similar
- Different forms of SQL are found in different DBMS's (Database Management Systems)

Database Management Systems

- MySQL, SQLite, PostgreSQL, Microsoft SQL Server are all popular DBMS's
- The different DBMS's are not very important right now
- Just know that the SQL syntax varies between the different DBMS's
- We will be using MySQL as an example

Basic SQL Table

```
CREATE TABLE people (  
  username VARCHAR(255),  
  first_name VARCHAR(255),  
  last_name VARCHAR(255),  
  age INT  
);
```

- Between the parenthesis the first value is the column name
- The second value is the **DATA TYPE**
- If the column should hold a number, you use the **INT** type
- If the column should hold a mixture of characters you use the **VARCHAR** type
- The VARCHAR type has a length parameter. We use 255 because it is suitable and common.
- The code between the parenthesis defines the **schema** for the table

Constraints

- PRIMARY KEY - Value(s) in specified column(s) must be unique for each row in a table and not be NULL
- UNIQUE - Value(s) in specified column(s) must be unique for each row in a table
- NOT NULL - Values for the column must not be NULL
- DEFAULT - Provide a default value for the column if no value is provided
- FOREIGN KEY - Value(s) in specified column(s) must reference an existing record in another table

Please **write these down**. You will need them for an exercise. We will discuss foreign key later.

SQL Table With Constraints

```
CREATE TABLE people (  
  username VARCHAR(255) NOT NULL UNIQUE,  
  first_name VARCHAR(255),  
  last_name VARCHAR(255) NOT NULL,  
  age INT DEFAULT 21  
);
```

- When a column is NOT NULL then a row cannot be inserted that has no value for that column
- For example, we don't want any user record in our database with no username
- Lack of a username could cause bugs or broken features in our codebase so we ensure that one exists for every user
- The DEFAULT constraint provides a default value for a column value when no value is provided
- The UNIQUE constraint ensures that no two records can share the same value for the given column

Constraints - Review

- Correct data is very important to a stable program
- Constraints on the DB level enforce good data
- Make your DB as strict as possible
- If your application expects a value to be unique, you must remember to create that constraint
- If a value should never be empty, you must enforce the NOT NULL constraint

PRIMARY KEY

Most tables include an ID column. It holds a unique value for each row in a table.

Most ID columns are automatically populated or set to "auto increment".

```
CREATE TABLE people (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY  
  username VARCHAR(255) NOT NULL UNIQUE,  
  first_name VARCHAR(255),  
  last_name VARCHAR(255) NOT NULL,  
  age INT DEFAULT 21  
);
```

Exercise - Create SQL Table Together

Lets create a users table with the following columns:

- `id`
- `name`
- `email`
- `password`
- `created_at`
- `updated_at`
- `num_logins`
- `is_email_address_confirmed`

Exercise - Create SQL Table Together

```
CREATE TABLE users (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  name VARCHAR(255) NOT NULL,  
  email VARCHAR(255) NOT NULL UNIQUE,  
  password VARCHAR(255) NOT NULL,  
  created_at DATETIME NOT NULL,  
  updated_at DATETIME NOT NULL,  
  num_logins INT NOT NULL DEFAULT 0,  
  is_email_address_confirmed TINYINT(1) NOT NULL /* boolean */  
);
```

- TINYINT is typically used for boolean types in MySQL
- Data types can vary between DBMS. PostgreSQL has a data type actually called "boolean"
- We don't want any duplicate email addresses in our database so we enforce that with the UNIQUE constraint

Exercise - Create Products Table

- Pair up and help each other create a products table
- Include the columns id, name, price, is_in_stock, sku
- SKU is unique ID for product, can contain letters and numbers
- Assign the correct data type to each column
- Assign a PRIMARY KEY (auto incrementing)
- Assign a UNIQUE KEY to the appropriate column
- Reference the users table for tips

```
CREATE TABLE users (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  name VARCHAR(255) NOT NULL,  
  email VARCHAR(255) NOT NULL UNIQUE,  
  password VARCHAR(255) NOT NULL,  
  created_at DATETIME NOT NULL,  
  updated_at DATETIME NOT NULL,  
  num_logins INT NOT NULL DEFAULT 0,  
  is_email_address_confirmed TINYINT(1) NOT NULL /* boolean */  
);
```

Exercise - Create Products Table - Answer

- Pair up and help each other create a products table
- Include the columns id, name, price, is_in_stock, sku
- SKU is unique ID for product, can contain letters and numbers
- Assign the correct data type to each column
- Assign a PRIMARY KEY (auto incrementing)
- Assign a UNIQUE KEY to the appropriate column
- Reference the users table for clues

```
CREATE TABLE products (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  name VARCHAR(255) NOT NULL,  
  price INT NOT NULL, /* Could also use DECIMAL */  
  is_in_stock TINYINT(1) NOT NULL, /* boolean */  
  sku VARCHAR(255) NOT NULL UNIQUE  
);
```

Relationships/Associations

- Imagine your application has users and users have profiles
- You store the user profile information in a separate table to reduce the number of columns in your users table ¹
- You need a way to associate a profile record with a user record
- A foreign key is used to create this relationship

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

¹ Imagine the profiles table has several other columns such as bio, location, etc.

Relationships/Associations

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

```
SELECT * FROM profiles WHERE profiles.user_id = 7;
```

id	user_id	photo
1	7	sally.jpg

Foreign Keys

You've seen that foreign keys can be used to associate data from one table to data from another table. Such as a specific profile record to a specific user record.

Foreign keys are more than just a column in one table referencing a row/record in another table.

Foreign keys enforce **constraints** and provide data integrity.

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

Foreign Keys

The foreign key constraint on the `user_id` column guarantees that no profile record can be inserted with a `user_id` value that does not exist. That would cause the database to contain invalid data.

- We can't have a profile tied to a non existing user
- This is referred to as a referential constraint
- Referential constraints keep the database consistent

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

Foreign Keys

Example constraint violation

This is what happens when you try to insert a profile record with a non existing user ID.

Cannot add or update a child row:

a foreign key constraint fails

```
(`nycda`.`profiles`, CONSTRAINT `profiles_ibfk_1` FOREIGN KEY (`user_id`) REFERENCES `users` (`id`))
```

Foreign Keys

- Referential constraints are even more important when it comes to deleting records
- If I delete Sally (user ID 7), is Sally's profile record still relevant?
The answer is probably no. Foreign key constraints can be configured to "cascade delete".
- You can have the profile record automatically deleted when the related user record is deleted
- The profile record is not used by any other table. It has a direct 1-to-1 relationship with a user.
(there is no point in keeping it)

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

Database Relationships: 1-to-1

- In a one-to-one relationship, a record in one table is related to only one record in another table
- Examples:
 - One `users` record has one `profiles` record
 - One `customers` record has one `customer_details` record
 - One `citizens` record has one `social_security_numbers` record

Database Relationships: 1-to-1

Implementing a 1-1 relationship

1. Add a column to one of the tables being linked together that **references** the **PRIMARY ID** of the other table
2. Add a **FOREIGN KEY CONSTRAINT** to the table with the column referencing the **PRIMARY ID** of the other table

In this case the `profiles` table will have a `user_id` column that references the `id` column on the `users` table.

+-----+			+-----+			
users			profiles			
+-----+			+-----+			
id	name		id	user_id	photo	
+-----+			+-----+			
7	Sally		1	7	sally.jpg	
8	Ned		2	9	ned.png	
9	Sally		3	8	s1.png	
+-----+			+-----+			

Database Relationships: 1-to-1

A true 1-to-1 relationship means that every user can have no more than one profile!

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

Quiz

What constraint would we use to ensure that the profiles table doesn't have more than one row/record with the same user_id value?

+-----+		
users		
+-----+		
id name		
+-----+		
7 Sally		
8 Ned		
9 Sally		
+-----+		

+-----+			
profiles			
+-----+			
id user_id photo			
+-----+			
1 7		sally.jpg	
2 9		ned.png	
3 8		s1.png	
+-----+			

Answer

A UNIQUE constraint

No 1-to-1 relationship is truly 1-to-1 without a UNIQUE constraint on the foreign key column.

```
CREATE TABLE users (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  name VARCHAR(255) NOT NULL  
);  
  
CREATE TABLE profiles (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  user_id INT(11) UNIQUE, /* References users.id UNIQUE */  
  photo VARCHAR(255),  
  FOREIGN KEY (user_id) REFERENCES users(id) ON DELETE CASCADE /* FOREIGN KEY CONSTRAINT */  
);
```

users			profiles		
id	name		id	user_id	photo
7	Sally		1	7	sally.jpg
8	Ned		2	9	ned.png
9	Sally		3	8	s1.png

Database Relationships: 1-to-many

- In a one-to-many relationship, a record in one table is related to many records in another table
- Examples:
 - One **users** record has many **blog_posts**
 - One **blog_posts** record has many **comments**
 - One **users** record has many **tweets**

Database Relationships: 1-to-many

- Users can have many blog posts
- Ned has two blog posts
- Sally has two blog posts

users			blog_posts		
id	name		id	user_id	title
7	Sally		1	7	foo
8	Ned		2	7	hello
9	Bob		3	8	bar
			4	8	baz

Database Relationships: 1-to-many

1-to-many is implemented the same way as a 1-to-1 but without the UNIQUE constraint on the FOREIGN KEY column.

1. Add a column to one of the tables being linked together that **references** the **PRIMARY ID** of the other table
2. Add a **FOREIGN KEY CONSTRAINT** to the table with the column referencing the **PRIMARY ID** of the other table

Database Relationships: 1-to-many

1. Add a column to one of the tables being linked together that **references** the **PRIMARY ID** of the other table
2. Add a **FOREIGN KEY CONSTRAINT** to the table with the column referencing the **PRIMARY ID** of the other table

```
CREATE TABLE users (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  name VARCHAR(255) NOT NULL  
);  
  
CREATE TABLE blog_posts (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  user_id INT(11) NOT NULL, /* References users.id NOT UNIQUE */  
  title VARCHAR(255),  
  FOREIGN KEY (user_id) REFERENCES users(id) ON DELETE CASCADE /* FOREIGN KEY CONSTRAINT */  
);
```

+-----+			+-----+			
users			blog_posts			
+-----+			+-----+			
id	name		id	user_id	title	
+-----+			+-----+			
7	Sally		1	7	foo	
8	Ned		2	7	hello	
9	Bob		3	8	bar	
+-----+			4	8	baz	
			+-----+			

Quiz

What is the difference between a 1-to-1 and 1-to-many relationship when it comes to CONSTRAINTS?

Answer

What is the difference between a 1-to-1 and 1-to-many relationship when it comes to CONSTRAINTS?

1-to-1 has a UNIQUE constraint on the column referencing the other table, 1-to-many does not

Structurally, there is no other difference.

Database Relationships: many-to-many

- In a many-to-many relationship, a record in one table is related to many records in another table and vice versa.
- Examples:
 - A student can take multiple classes
 - A class can have multiple students
 - Orders can have many Products
 - Products can belong to many Orders

Database Relationships: many-to-many

Implementing a many-to-many relationship involves a join table. Also known as a pivot table.

The pivot table will associate records together from two separate tables.

students			classes		
id	name		id	title	
7	Sally		1	Intro to Ruby	
8	Ned		2	WDI	
9	Bob		3	Javascript 101	

PIVOT TABLE

enrollments				
id	student_id	class_id		
7	7	1		
8	7	2		
9	8	2		

Database Relationships: many-to-many

Implementing a many-to-many relationship

1. Create a pivot table
2. Add a column for each associated table that will reference the PRIMARY KEY of the associated table
3. Add a FOREIGN KEY constraint for each column referencing a PRIMARY KEY of another table

```
CREATE TABLE enrollments (  
  id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  student_id INT(11) NOT NULL,  
  class_id INT(11) NOT NULL,  
  FOREIGN KEY (student_id) REFERENCES students(id) ON DELETE CASCADE,  
  FOREIGN KEY (class_id) REFERENCES classes(id) ON DELETE CASCADE  
);
```

students			classes		
id	name		id	title	
7	Sally		1	Intro to Ruby	
8	Ned		2	WDI	
9	Bob		3	Javascript 101	

PIVOT TABLE

enrollments				
id	student_id		class_id	
7	7		1	
8	7		2	
9	8		2	

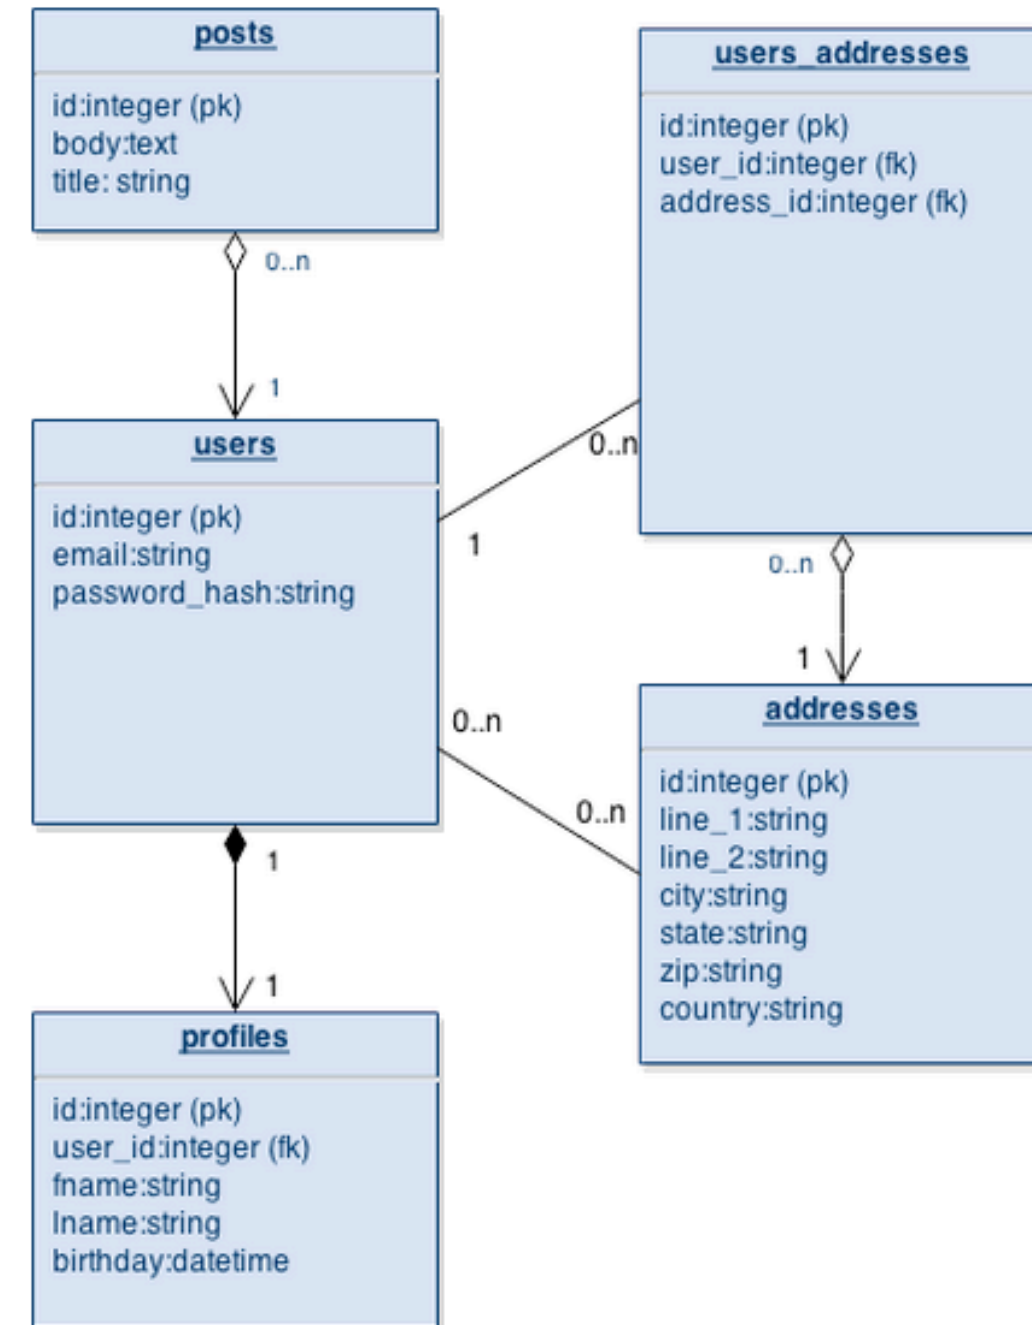
A full database diagram

- A "full" db diagram has all of your tables with their columns and data types
- The relationships between tables are illustrated with arrows

1 ---> 0..n = *one to many*

1 ---> 1 = *one to one*

0..n ---> 0..n = *many to many*



Exercise

- Let's build a database diagram for an existing major website together!

Exercise

- Develop a database diagram for one of your favorite websites.
- Start with the users table and go from there!
- Be sure to indicate the relationships between different tables in the database
- You could use www.draw.io or a pen and paper

A final word/summary

A **database** is composed of **tables**. Each **table** has **columns** and can be **related** to other tables in the database. Each column has a **data type**. Each table also has **rows** of data which correspond to the **column names** and data types described in the **schema**.

And they all lived happily ever after!

Resources

TeamTreeHouse

Database Foundations - Introduction to Data, Databases and SQL